



ABRAHAM OF LONDON — INSTITUTIONAL MANUAL

Version: 2.0 **Date:** April 2026 **Classification:** Internal — Institutional **Status:** Governing Document

Table of Contents

- [PART I — FOUNDING PRINCIPLES & DOCTRINE](#)
 - [Chapter 1: The Founding Problem](#)
 - [Chapter 2: The Institutional Thesis](#)
 - [Chapter 3: Product Philosophy](#)
 - [Chapter 4: Operating Principle](#)
 - [Chapter 5: Core Claims About Execution](#)
- [PART II — PRODUCT SYSTEMS](#)
 - [Chapter 6: Fast Diagnostic](#)
 - [Chapter 7: Purpose Alignment Assessment](#)
 - [Chapter 8: Constitutional Assessment](#)
 - [Chapter 9: The Contradiction Model](#)
 - [Chapter 10: Economic Framing](#)
 - [Chapter 11: Strategy Room](#)
 - [Chapter 12: Objection Handling](#)
 - [Chapter 13: Outcome Verification](#)
 - [Chapter 14: Boardroom Mode](#)
 - [Chapter 15: The Intelligence Spine](#)
- [PART III — PRODUCT ELEVATION LAYER](#)
 - [Chapter 16: Decision Authority Index v1](#)
 - [Chapter 17: Cost of Inaction Engine](#)
 - [Chapter 18: Execution Failure Layer](#)
 - [Chapter 19: Decision Memory v1](#)

- [PART IV — SALES & POSITIONING](#)
 - [Chapter 20: Sales Doctrine](#)
 - [Chapter 21: Enterprise Case System](#)
 - [PART V — CONTENT & DISTRIBUTION](#)
 - [Chapter 22: Content Amplification](#)
 - [Chapter 23: Engagement Lane Architecture](#)
 - [PART VI — SYSTEM ARCHITECTURE](#)
 - [Chapter 24: The 13-Layer Technical Architecture](#)
 - [Chapter 25: Delivery Model & Operating Cadence](#)
 - [PART VII — CAPITAL & INVESTOR RELATIONS](#)
 - [Chapter 26: Capital Narrative](#)
 - [Chapter 27: Investor Memo](#)
 - [Chapter 28: Investor Call Script](#)
 - [Chapter 29: Objection Handling — Investor Warbook](#)
 - [Chapter 30: Fundraising Strategy](#)
 - [PART VIII — OPERATIONAL DISCIPLINE & GOVERNANCE](#)
 - [Chapter 31: The First Ten Accounts](#)
 - [Chapter 32: What Not To Do](#)
 - [Chapter 33: Quality Assurance & Compliance](#)
 - [Chapter 34: Engagement Termination](#)
 - [PART IX — IP PROTECTION & COMPETITIVE DEFENSE](#)
 - [Chapter 35: Disclosure Principle](#)
 - [Chapter 36: Public vs Proprietary Layers](#)
 - [APPENDICES](#)
 - [Appendix A: Pricing Matrix](#)
 - [Appendix B: Severity Taxonomy](#)
 - [Appendix C: Condition Class Taxonomy](#)
 - [Appendix D: 13 Contradiction Archetype Registry](#)
 - [Appendix E: Quality Gate Checklist](#)
-

PART I — FOUNDING PRINCIPLES & DOCTRINE

Chapter 1: The Founding Problem

Organisations do not fail because they lack intelligence. They fail because they cannot resolve decisions.

Every organisation has access to advisors, consultants, data dashboards, and strategic frameworks. The modern enterprise is saturated with information. And yet decisions stall. They stall not because the answer is unknown, but because no one has the authority, the clarity, or the institutional pressure to force the decision to conclusion.

The founding problem that Abraham of London addresses is structural, not informational. It is the observation that the gap between knowing what to do and actually doing it is not a knowledge gap — it is a decision authority gap. Someone knows the answer. No one has the mandate to act on it. Or the mandate exists on paper but has been quietly hollowed out by politics, ambiguity, or avoidance.

This problem compounds. An unresolved decision does not sit still. It degrades. Options compress. The cost of delay increases. The people who should be deciding begin to act from different assumptions. What was a clear choice at Week 1 becomes a structural emergency by Week 12 — not because anything changed externally, but because the internal decision architecture failed to force resolution.

The conventional response to this is more information: another report, another workshop, another alignment session. These interventions treat the symptom. They add intelligence to a system that already has enough intelligence. They do not address the structural condition that prevents the intelligence from being acted upon.

Abraham of London was built on a different premise: the problem is not what the organisation knows. The problem is what the organisation refuses to decide. The system exists to identify the contradiction, price the consequence of inaction, force the decision to a named owner, and verify that execution actually occurs. It is not advisory. It is infrastructure.

This distinction matters. Advisory depends on the client's willingness to listen. Infrastructure operates regardless. The system does not require agreement. It requires accuracy. If the diagnosis is correct, the pressure is legitimate. If the pressure is legitimate, the decision eventually resolves — either by the client's action or by the consequences the system predicted arriving on schedule.

The founding problem is therefore not "organisations need better advice." The founding problem is: organisations lack a structural mechanism that identifies unresolved decisions, assigns consequence to inaction, and enforces resolution. That mechanism is what this institution builds.

Chapter 2: The Institutional Thesis

The institutional thesis of Abraham of London can be stated in four operations:

1. **Identify the contradiction.** Every stalled decision contains a contradiction — a gap between what the organisation says it wants and what it is actually doing. The system surfaces this contradiction through structured diagnostic questioning, deterministic classification, and cross-session pattern detection. The contradiction is not invented. It is derived from the client's own statements, scored for fidelity, and classified into a condition class.
2. **Price the consequence.** Once the contradiction is identified, the system computes what it costs to leave it unresolved. This is not a fabricated number. When the user provides financial anchors, the system uses them. When they do not, the system describes the structural exposure in qualitative terms: what happens at 30 days, 60 days, 90 days. The pricing of consequence converts an abstract problem into a concrete liability. It makes inaction expensive in the client's own terms.
3. **Force the decision.** The system does not suggest. It forces. Every diagnostic culminates in a named owner, a specific action, and a deadline. The Pre-Commitment Gate in the Fast Diagnostic asks the user whether they are willing to act within 48 hours. The Strategy Room locks a decision, names an owner, sequences moves, and tracks execution. The system creates institutional pressure where none existed before.
4. **Verify the action.** Diagnosis without follow-through is entertainment. The system tracks whether the committed action was taken, whether the structural condition changed, and whether the outcome was real or cosmetic. The Outcome Verification loop measures actual change — structural change, temporary fix, or no change — and feeds this back into the Intelligence Spine for future sessions.

These four operations are not features. They are the thesis. Every product, every surface, every pricing decision, and every engagement protocol derives from this sequence. If a proposed feature does not serve one of these four operations, it does not belong in the system.

The thesis also implies a market position: Abraham of London is not a consulting firm. It is not a coaching practice. It is not a SaaS analytics tool. It is Decision Authority Infrastructure — a new category. The system creates the structural conditions under which decisions get made. It does not depend on the client's goodwill, motivation, or self-awareness. It depends on accuracy. If the system is accurate, the client either acts or watches the predicted consequences arrive.

Chapter 3: Product Philosophy

The product philosophy of Abraham of London is governed by four principles:

Forceful diagnosis. The system does not ask what the client wants to hear. It asks what exposes the truth. Questions are designed to surface contradictions, not to validate existing narratives. The Fast Diagnostic does not ask "What are your goals?" It asks "What decision are you avoiding?" and "If you had to decide in 24 hours, what would you do?" These questions are uncomfortable by design. Comfort produces vague input. Vague input produces vague output. The system requires specificity to function. The C3 Fidelity Scorer (Clarity, Context, Consequence) enforces this: if input lacks specificity, the system enters recovery mode rather than producing a diluted diagnosis.

Enforcement over suggestion. The system does not make recommendations and hope they are followed. It creates structural enforcement. The Pre-Commitment Gate captures willingness to act. The Strategy Room locks decisions with named owners and deadlines. The Escalation Engine applies time-based pressure at 24h, 48h, 72h, 5d, and 7d intervals. The Breach Ladder detects when commitments are violated. These are not features layered on top of an advisory model. They are the model. The system assumes that without enforcement, decisions revert to their default path — which is almost always inaction.

Pressure as mechanism. Pressure is not punitive. It is structural. The system uses pressure the way a combustion engine uses compression: to force action from potential energy. The pressureIndex on the Intelligence Spine is a monotonically increasing value (0-100). It never decreases. Each stage, each missed deadline, each detected contradiction increases pressure. This is not gamification. It is a structural representation of the fact that unresolved decisions accumulate consequence over time.

Precision over comfort. The system would rather say nothing than say something inaccurate. The Arbiter Tournament runs five mandatory validation rules against every synthesis output. If any hard rule fails, the synthesis is rejected and the user sees the mismatch — not a quietly degraded result. The system does not bluff intelligence. If the C3 score is below threshold, the system enters HARD_RECOVERY or SOFT_RECOVERY mode rather than fabricating depth from shallow input. This precision comes at a cost: some users will feel the system is too harsh. Those users are not the target market.

Chapter 4: Operating Principle

Authority derives from accuracy, not permission.

This is the operating principle of the institution. It means that the system does not ask for the client's approval before delivering a diagnosis. It does not soften findings to maintain the relationship. It does not adjust its output based on what the client wants to hear. The system derives its authority from one thing: being right.

If the system identifies a contradiction between what the client says and what the client does, and that contradiction is accurate, the system has earned the authority to press on it. If the system predicts that inaction will produce a specific structural consequence within 90 days, and that prediction is accurate, the system has earned the authority to escalate.

This principle has a corollary: the system repels wrong clients. Not every organisation is ready for this kind of intervention. Some want validation. Some want a softer version of what they already believe. Some want to appear decisive without actually deciding. The system is designed to make these clients uncomfortable — not through hostility, but through accuracy. The Do-Not-Sell flag exists for this reason. When the system detects that engagement would be unproductive — because the client lacks authority to act, because the problem is not real, or because the client is seeking comfort rather than resolution — the system withdraws the offer rather than diluting the product.

The False Authority Detection mechanism serves the same purpose. When a user claims ownership of a decision but their answers reveal they cannot actually force it to conclusion, the system flags the discrepancy. It does not pretend the user has authority they do not have. It names the gap. This is uncomfortable. It is also accurate. And accuracy is the source of authority.

The system does not compete on warmth, accessibility, or ease of use. It competes on being right. Every engagement protocol, every quality gate, every validation rule exists to protect this single asset. If the system is accurate, clients return. If it is inaccurate, no amount of branding will save it.

Operational implications: no feature ships that trades accuracy for user comfort. No output reaches the client that has not passed the Quality Gate. No synthesis is delivered that the Arbiter Tournament has not validated. No economic projection is made without a user-provided anchor. The system's credibility is its only moat, and every decision in the codebase protects it.

Chapter 5: Core Claims About Execution

The system makes five core claims about what it does. Each claim is verifiable against the codebase and must remain accurate. If any claim becomes false, the system is in breach of its own doctrine.

Claim 1: The system tracks decisions. Every user interaction begins with a decision stated in the user's own words. This decision is captured in the CaseObject, classified into a condition class (authority, definition, execution, instability), scored for specificity via the C3 Fidelity Scorer, and persisted across sessions through the Intelligence Spine. The system does not lose context between stages. The CaseObject created in the Fast Diagnostic is the same CaseObject consumed by the Constitutional Assessment, the Strategy Room, and Executive Reporting. No stage re-derives context from zero.

Claim 2: The system enforces commitments. The Pre-Commitment Gate in the Fast Diagnostic asks whether the user is willing to act within 48 hours. This intent is captured on the spine (`preCommitment.willing48h`). The Strategy Room locks a decision with a named owner, sequenced moves, and execution tracking. The Escalation Engine applies time-based pressure at defined intervals (24h, 48h, 72h, 5d, 7d). These are not suggestions. They are structural enforcement mechanisms that create accountability where none existed.

Claim 3: The system escalates breaches. When a committed action is not taken, the system detects the breach via the execution tracking on the spine (`execution.breach` , `execution.breachCount`). The Breach Ladder applies graduated escalation. The system does not silently accept missed deadlines. It increases pressure, re-surfaces the contradiction, and records the breach in the Evidence Ledger. The `pressureIndex` increases. The `integrityScore` may decrease if gaming behaviour is detected.

Claim 4: The system reports contradictions. The deterministic output on every spine includes a `contradictionSet` — an array of specific contradictions derived from the user's own answers. These are not generic observations. They are structural mismatches between what the user said in one question and what they said in another. The synthesis layer (when activated) produces a `primaryContradiction` that references the deterministic set. If the synthesis contradicts the deterministic finding, the Arbiter Tournament rejects it.

Claim 5: The system verifies outcomes. The Outcome Verification loop asks whether the action was taken, what impact it produced, and whether the structural condition changed. The `verifiedImpact` field on the spine records one of three values: `structural_change` , `temporary_fix` , or `no_change` . This feedback is appended to the Evidence Ledger and informs future sessions through the Memory system. The system does not assume its own effectiveness. It measures it.

PART II — PRODUCT SYSTEMS

Chapter 6: Fast Diagnostic

The Fast Diagnostic is the entry-level product. It is free. It takes under two minutes. All scoring is server-side — no client-side computation occurs.

UX Flow:

1. **Q1 — The Decision.** The user states the decision they are facing in their own words. This becomes the `decision` field on the CaseObject. The system does not provide options or categories. The user must articulate the problem.
2. **Q2 — Structured Ownership.** The user identifies who owns the decision. This becomes the `claimedOwner` field. The system later compares this against inferred ownership signals to detect authority gaps.
3. **Initial Signal Block (ownership-aware).** Based on Q1 and Q2, the system generates an initial signal block that is aware of the stated ownership structure. This is the first diagnostic output the user sees. It is deterministic and derived from the CaseObject without LLM involvement.
4. **Commitment Gate.** Before proceeding, the user is asked whether they are willing to act on the diagnosis within 48 hours. This is captured as `preCommitment.willing48h` on the Intelligence Spine. The gate is not optional. It separates those who want information from those who intend to act.
5. **Q3-Q6 — Deepening Questions.** These questions probe prior attempts (`priorAttempt`), cost of delay (`costOfDelay`), blockers (`blocker`), and forced action under constraint (`forcedAction`). Each answer feeds the CaseObject and improves the C3 specificity score.
6. **Server-Side Scoring.** The complete CaseObject is scored server-side. The C3 Fidelity Scorer computes Clarity, Context, and Consequence scores (each 0-1). The composite specificity score determines the tier: below 0.5 enters `HARD_RECOVERY` (no synthesis, interrogation only), 0.5-0.7 enters `SOFT_RECOVERY` (deterministic only), 0.7 and above enters `FULL_SYNTHESIS` (full LLM synthesis plus Arbiter Tournament validation).
7. **Result with ExecutiveDecisionAuthorityBlock.** The final output includes the deterministic classification (condition class, signal, contradiction set, blocker class), the Default Path Forecast (what happens at 7/30/90 days if they do not act), and — if the C3 tier permits — a governed synthesis with primary contradiction, concrete move, and avoidance surface. The result includes the Pre-Commitment status and a clear next action.

The Fast Diagnostic creates the Intelligence Spine. Every subsequent product reads from and writes to this spine. The diagnostic is not a standalone tool. It is the entry point to a persistent decision-tracking infrastructure.

Chapter 7: Purpose Alignment Assessment

The Purpose Alignment Assessment is a personal diagnostic product. It is free.

This assessment operates on a dual-axis model: **Resonance** and **Certainty**. Resonance measures the degree to which a person's stated direction aligns with their observable behaviour and emotional responses. Certainty measures the degree to which they are clear about what they actually want, as opposed to what they believe they should want.

The dual-axis model produces four quadrants:

- **High Resonance / High Certainty:** Aligned. The person knows what they want and their actions reflect it.
- **High Resonance / Low Certainty:** Instinctive alignment without articulated direction. The person is moving in the right direction but cannot explain why.
- **Low Resonance / High Certainty:** Stated clarity but structural misalignment. The person can articulate what they want but their behaviour contradicts it.
- **Low Resonance / Low Certainty:** Disordered. Neither direction nor alignment is present.

The system classifies results into alignment bands: **ALIGNED**, **DRIFTING**, **MISALIGNED**, and **DISORDERED**. These bands map to the `AlignmentBand` enum in the database schema. The assessment covers six domains: **IDENTITY**, **DECISION**, **ENVIRONMENT**, **BEHAVIOUR**, **EMOTIONAL_ORDER**, and **LEGACY**.

Pattern detection identifies recurring themes across responses. When a user completes the assessment multiple times, the system detects whether their alignment is improving, stable, or degrading across sessions. This longitudinal signal feeds into the Decision Memory system and can inform subsequent diagnostic interactions.

The Purpose Alignment Assessment serves a specific function in the product ladder: it establishes personal mandate clarity before the system applies institutional pressure. A user who lacks personal clarity will not benefit from the Strategy Room. The system uses Purpose Alignment signals to restrict escalation when a mandate vacuum is detected (see Chapter 6, Decision Directive system).

The assessment produces a report with domain-level scores, an overall alignment band, and specific correction recommendations. Reports can be exported as PDF. Enterprise campaigns can deploy the assessment across teams to detect collective alignment patterns.

Chapter 8: Constitutional Assessment

The Constitutional Assessment is a structural diagnostic. It consists of 10 questions and produces 9 domain scores. It is the system's primary instrument for evaluating organisational decision-making health.

The 9 domains assessed are:

1. **Decision Clarity** — whether decisions are defined with enough specificity to act on.
2. **Ownership Structure** — whether someone has clear, undisputed authority to decide.
3. **Execution Readiness** — whether the organisation can translate a decision into action.
4. **Accountability Architecture** — whether there are consequences for inaction.
5. **Communication Integrity** — whether what is said internally matches what is done.
6. **Strategic Coherence** — whether individual decisions align with stated direction.
7. **Resource Alignment** — whether resources are allocated to match stated priorities.
8. **Conflict Resolution** — whether disagreements are resolved or left to fester.
9. **Temporal Discipline** — whether decisions are made within appropriate timeframes.

Each domain is scored on a 1-5 scale per question. The assessment produces a total score, a percentage, a severity classification (low, moderate, high, critical), and a score band (stable, watch, fragile, escalate).

The severity classification uses the following thresholds: 80+ is "low" severity (stable with minor correction points), 60-79 is "moderate" (recoverable but drift is present), 40-59 is "high" (material weakness detected), below 40 is "critical" (immediate correction required).

The Constitutional Assessment can be deployed as a team campaign, where multiple respondents answer independently and the system aggregates results. Cross-respondent analysis detects where different stakeholders see the same domain differently — a signal of structural misalignment that no individual respondent would surface alone.

Results feed into the Intelligence Spine. The assessment deepens the CaseObject created by the Fast Diagnostic, adding domain-level granularity to the initial classification. The Constitutional Assessment is not a replacement for the Fast Diagnostic. It is the next layer of structural analysis, available after the entry-level diagnostic has established the case.

Chapter 9: The Contradiction Model

The contradiction model is the intellectual core of the system. Every stalled decision contains at least one contradiction — a structural mismatch between what the organisation says and what it does.

The system currently implements **4 core condition classes** with variant expressions:

1. **Authority** — the decision has no legitimate owner, or the stated owner lacks actual authority. Variant: someone claims ownership but cannot compel action.
2. **Definition** — the decision itself is not clearly defined. People are arguing about what to do before agreeing on what the question is. Variant: the decision has been reframed so many times that its original urgency has been diluted.
3. **Execution** — the decision has been made but is not being executed. The blocker is not informational but structural — resources, politics, or avoidance. Variant: partial execution creates the appearance of progress while the core issue remains unresolved.
4. **Instability** — the conditions around the decision are shifting faster than the organisation can adapt. The decision may be correct today and obsolete next week. Variant: the organisation is making stable decisions in an unstable environment, or unstable decisions in a stable one.

TARGET — not yet implemented: The system targets a taxonomy of **13 specific contradiction archetypes** distributed across these 4 condition classes. The 13-archetype registry is documented in Appendix D. The current implementation uses the 4 condition classes with the `detectContradictions` function performing rule-based contradiction detection (urgency without timeline, growth without commercial signals, strategic claims without depth, transformation without execution approach). The 13-archetype taxonomy represents the target state where each condition class contains 3-4 named archetypes with specific detection signatures, scoring weights, and prescribed interventions.

The contradiction model operates deterministically. Contradictions are derived from the user's own answers, not from statistical inference or sentiment analysis. The CaseObject stores the `contradiction` field (the primary contradiction identified), and the DeterministicOutput on the Intelligence Spine contains a `contradictionSet` — an array of all detected contradictions. The Arbiter Tournament validates that any LLM synthesis references these deterministic contradictions rather than inventing its own.

This model is the foundation of the system's authority. When the system tells a client "You say you want X, but your answers reveal you are doing Y," the authority of that statement comes from the client's own words, not from the

system's opinion. The contradiction is provable. That provability is what separates this system from advisory.

Chapter 10: Economic Framing

The economic framing of the system is handled by the Cost of Delay Engine, which operates server-side only. No internal formulas, multipliers, or scoring thresholds are exposed to the client.

The system produces economic pressure through two mechanisms:

Exposure Bands. Every case is classified into one of five exposure bands: low, moderate, high, critical, or undisclosed. The "undisclosed" band applies when the user has not provided financial anchors — the system does not fabricate financial exposure from qualitative input alone. When financial anchors are provided (`estimatedExposureGBP` , `headcountAffected` , `decisionValue`), the system uses them to calibrate the band. When they are not, the system derives the band from the organisational state (ORDERED, DRIFTING, MISALIGNED, DISORDERED).

Horizon Narratives. The system produces qualitative projections at 30-day, 60-day, and 90-day horizons. These narratives describe what happens if the decision remains unresolved. They vary by state, not by formula. For a DISORDERED state: "Control begins shifting to whoever acts first" (7 days), "The decision becomes harder to recover without escalation" (30 days), "The unresolved condition becomes part of normal operating behaviour" (90 days). For a DRIFTING state: "Execution slows because ownership is not settled" (7 days), "Stakeholders begin acting from different assumptions" (30 days).

The economic framing is financial only when user-anchored. If the user provides a revenue exposure figure or a headcount number, the system incorporates it. If they do not, the system speaks in structural terms — control transfer, option compression, decision decay — rather than inventing pound figures.

The Default Path Forecast on the Intelligence Spine adds mathematical rigour: `optionDecayRate` (options lost per month, 0-1 scale), `controlShiftProbability` (probability that decision control transfers away from stated owner, 0-1), and `structuralRiskShift` (stable, accelerating, or compounding). These are computed server-side from condition class base rates and user-provided specificity.

The economic framing serves the thesis: it prices the consequence of inaction. It does not advise. It quantifies what doing nothing costs, in terms the client has already anchored.

Chapter 11: Strategy Room

The Strategy Room is the enforcement product. It converts a diagnosed condition into a locked decision with a named owner, sequenced moves, and execution tracking. Pricing: **£750-£1,250**.

The Strategy Room operates on a 48-hour action-forcing model. Upon entry, the user's Intelligence Spine is consumed and the session is initialised with the accumulated diagnostic intelligence. The room does not start from zero. It starts from everything the system already knows about the case.

Core mechanics:

- **Locked Decision.** The decision is stated, confirmed, and locked. It cannot be redefined mid-session without the system flagging the redefinition as a potential avoidance signal.

- **Named Owner.** A specific person is identified as the decision owner. Not a role. Not a committee. A named individual with the authority to act.
- **Sequenced Moves.** The system produces a sequence of concrete actions — not goals, not aspirations, but specific moves with deadlines. Each move is tracked independently.
- **Execution Tracking.** Post-session, the system tracks whether each move was executed. The `execution` object on the Intelligence Spine records breaches, action timestamps, blocker reports, and verified impact.

The Strategy Room is gated. Not every user qualifies. The Decision Directive system evaluates whether the user's diagnostic state supports strategic intervention. If structural failure or unmanaged risk is detected, escalation to the Strategy Room is blocked. If a mandate vacuum exists, escalation is restricted. The system does not sell strategic intervention to users whose structural condition cannot support it.

Session initialisation, follow-up sequences, and conversion tracking are managed through dedicated API routes. The Strategy Room persists enforcement data through Prisma models including session state, execution milestones, and breach records (database migrations: `20260420_add_strategy_room_tables` , `20260423_add_strategy_room_enforcement_persistence`).

The Strategy Room is where the system's thesis is most visible: force the decision, name the owner, track the execution, escalate the breach. It is not a workshop. It is an enforcement mechanism with a 48-hour clock.

Chapter 12: Objection Handling

The system maintains a registry of 20 objection patterns with scripted responses. These objections cover the most common resistance patterns encountered during sales, onboarding, and engagement.

The objections fall into four categories:

Authority objections (5 patterns): "I need to check with my board," "This isn't my decision to make," "We need consensus first," "Our governance structure doesn't work this way," "I don't have budget authority." The response pattern for all authority objections is the same: surface the contradiction between the stated urgency and the claimed inability to act. If it is urgent enough to investigate, it is urgent enough to decide.

Value objections (5 patterns): "This is too expensive," "We can do this internally," "We already have consultants," "What's the ROI?", "We need to see a demo first." The response pattern: reframe from cost to cost of inaction. The question is not "What does this cost?" but "What does not doing this cost you per month?"

Timing objections (5 patterns): "Not the right time," "We're in the middle of a restructure," "Let's revisit next quarter," "We need to finish our current initiative first," "Things are too chaotic right now." The response pattern: timing objections are often the strongest indicator that the intervention is needed now. Chaos is not a reason to delay decision clarity. It is the reason decision clarity is urgent.

Legitimacy objections (5 patterns): "How is this different from consulting?", "What credentials do you have?", "Can you show me case studies?", "This sounds like coaching," "Why should I trust a system over my own judgement?" The response pattern: the system does not ask for trust. It asks for accuracy verification. Run the diagnostic. If the findings are accurate, the system has earned authority. If they are not, the engagement should not proceed.

Each objection pattern includes: the objection text, the underlying belief, the structural response, and a closing question that re-establishes engagement. The scripts are used in sales calls, follow-up sequences, and automated engagement flows. They are not persuasion techniques. They are applications of the system's own doctrine to the sales surface.

Chapter 13: Outcome Verification

Outcome Verification closes the loop. Without it, the system is advisory — it diagnoses and recommends, but never verifies whether anything changed. With it, the system is infrastructure — it measures its own impact and uses that measurement to improve.

The Outcome Verification loop captures three signals:

1. **Action Taken.** Was the committed action executed? This is a binary recorded on the spine as `execution.actionTaken` with a timestamp `execution.actionTakenAt`. If the action was not taken, the system records a breach and the Escalation Engine applies pressure.
2. **Blocker Reported.** If the action was not taken, the user can report what blocked execution. This is captured as `execution.blockerReported`. The system uses this to distinguish between avoidance (the user chose not to act) and legitimate obstruction (something external prevented action). The distinction matters for future diagnosis.
3. **Verified Impact.** When the action is taken, the system asks whether it produced structural change, a temporary fix, or no change. This is recorded as `execution.verifiedImpact` with three possible values: `structural_change`, `temporary_fix`, `no_change`. This is the most important signal in the system. It is the difference between "we did something" and "something actually changed."

The verified impact feeds back into the Intelligence Spine and is appended to the Evidence Ledger. Over time, this creates a longitudinal record of what actually works for this client, this type of decision, and this condition class. The Outcome Verification data also feeds the accuracy feedback loop — the user can confirm whether the system's diagnosis was accurate (`accuracyFeedback.response` : "yes", "partial", "no").

The outcome feedback loop is what allows the system to make its fifth core claim: it verifies outcomes. Diagnosis without verification is speculation. The system does not speculate. It measures.

Chapter 14: Boardroom Mode

Boardroom Mode is a multi-actor deployment that produces a Board Dossier — a decision document designed to survive CFO scrutiny, board challenge, and political resistance. It answers three questions: "Why this?", "Why now?", and "What happens if we don't?"

Activation Gate. Boardroom Mode is not available by default. It activates when the estimated monthly cost is at or above £5,000 and the accuracy feedback is "yes" or "partial", or when the monthly cost is at or above £20,000 regardless of accuracy feedback. Below these thresholds, the system returns: "This is not a board-level issue. Resolve operationally." The classification is `BOARD_RESTRICTED`.

Board Dossier Structure. The dossier contains:

- Sections with three tonal modes: factual, confrontational, and quantified. Each section is generated from the Intelligence Spine.
- Objection handling: pre-built responses to anticipated board objections, specific to the case.
- Decision path: two or three options with consequences for each, and a recommended path.

Multi-Stakeholder Collision Detection. When Boardroom Mode is deployed across multiple stakeholders, the system detects where their answers diverge. This is collision detection — not consensus-building. If the CEO says the problem is execution and the COO says the problem is authority, the system does not average their views. It surfaces the collision and identifies what it means for the decision.

The collision detection operates through the StakeholderMap on the Intelligence Spine, which tracks `formalOwner`, `realOwner`, `blockers`, `silentInfluencers`, and `misalignedParties`. When multiple stakeholders provide input, the system identifies authority gaps — situations where no one has clear authority — and cost divergence — situations where stakeholders assign radically different values to the same decision.

Boardroom Mode is where the system operates at its highest institutional level. It produces decision-grade documents that can be tabled in a board meeting, not as a consultant's recommendation but as a structural analysis of the decision condition backed by the client's own data.

Chapter 15: The Intelligence Spine

The Intelligence Spine is the single contract across all stages. It is the most important technical construct in the system.

Every stage reads from the spine. Every stage writes to the spine. No stage computes in isolation. No stage feels like a reset. The Fast Diagnostic creates it. The Constitutional Assessment deepens it. The Team Assessment validates it. Enterprise prices it. Executive Reporting enforces it. The Strategy Room simulates it. Outcome Verification verifies it.

Spine Components:

- **CaseObject** — the user's decision in their own words, with structured fields for ownership, blockers, cost of delay, and forced action scenarios.
- **C3Score** — the Clarity/Context/Consequence fidelity score with tiered enforcement (`HARD_RECOVERY < 0.5`, `SOFT_RECOVERY 0.5-0.7`, `FULL_SYNTHESIS >= 0.7`) and confidence band (low/medium/high).
- **DeterministicOutput** — the non-negotiable classification: condition class, signal, contradiction set, blocker class. Always computed. Never nullable.
- **GovernedSynthesis** — the LLM-generated analysis, null if the C3 tier is below `FULL_SYNTHESIS` or if the Arbiter Tournament rejected it.
- **DefaultPathForecast** — what happens if they do not act, with mathematical derivations for option decay rate, control shift probability, and structural risk shift.
- **Memory** — cross-session recurrence detection. Null for first-time users.
- **StakeholderMap** — formal owner, real owner, blockers, silent influencers, misaligned parties. Null until populated by multi-actor deployment.
- **Stage** — current position in the ladder (`fast_diagnostic`, `constitutional`, `team`, `enterprise`, `executive_reporting`, `strategy_room`, `outcome_verification`).
- **History** — accumulated event log from every completed stage, with timestamps, snapshots, and contribution summaries.
- **preCommitment** — the user's willingness to act within 48 hours, captured at the Commitment Gate.
- **accuracyFeedback** — the user's response on whether the diagnosis was accurate (yes/partial/no).
- **economics** — user-stated cost data: estimated monthly cost, cost of delay monthly, decision owner, deadline.

- **flags** — system signals: avoidance suspected, economic sanity suspicious, do-not-sell triggered, false authority detected.
- **execution** — post-purchase tracking: breach status, action taken, blocker reported, verified impact.
- **integrityScore** — anti-gaming score (0-1). Below 0.5 indicates degraded integrity.
- **pressureIndex** — cross-stage pressure (0-100). Must only increase. Never resets.

The spine never resets between stages. This is a design invariant. Every user starts at the Fast Diagnostic and accumulates intelligence through every subsequent interaction. The spine is the institutional memory of the engagement.

PART III — PRODUCT ELEVATION LAYER

Chapter 16: Decision Authority Index v1

The Decision Authority Index (DAI) v1 is an interpretive layer that translates the system's raw diagnostic output into board-readable governance language. It does not expose numeric scores. It produces interpretive bands with governance meaning.

Interpretive Bands:

- **Strong** — decision authority is concentrated, exercised, and producing results. The organisation can decide and execute. Governance posture: maintain current authority structure with periodic review.
- **Strained** — decision authority exists but is under pressure. Execution is delayed or partial. One or more condition classes show early signs of drift. Governance posture: identify the strain source and resolve within current cycle.
- **Weak** — decision authority is fragmented, contested, or nominally assigned to someone who cannot exercise it. Decisions are being made by default rather than by design. Governance posture: restructure authority before attempting strategic initiatives.
- **Critical** — decision authority has collapsed. No one is effectively deciding. The organisation is operating on momentum, habit, or crisis response. Governance posture: emergency intervention. Stop new initiatives until authority structure is restored.

Board Meaning. Each band carries a specific governance implication. The DAI does not tell the board what to do. It tells the board what condition their decision-making infrastructure is in. This is a diagnostic, not a recommendation. The board decides what to do about the condition.

Next Governance Move. For each band, the system recommends the next governance-level action — not an operational action, but a governance action. For "Strained": convene a decision authority review within 14 days. For "Weak": suspend new strategic commitments until authority structure is clarified. For "Critical": appoint an interim decision authority with explicit mandate and timeline.

The DAI v1 is derived from the Intelligence Spine's accumulated data: condition class, contradiction count, breach history, pressure index, integrity score, and accuracy feedback. The derivation logic is server-side only. The output is qualitative. No numeric scores are exposed to the client.

Chapter 17: Cost of Inaction Engine

The Cost of Inaction (COI) Engine is a server-only computation layer that produces qualitative exposure assessments and horizon projections. It is implemented in `lib/server/decision/cost-of-inaction.server.ts` with a `server-only` import guard that prevents client-side execution.

Input Parameters:

- Organisational state: ORDERED, DRIFTING, MISALIGNED, or DISORDERED
- Estimated financial exposure in GBP (optional)
- Decision window (optional)
- Headcount affected (optional)
- Market exposure context (optional)

Output:

- `exposureBand` : low, moderate, high, critical, or undisclosed
- `horizon30` : qualitative narrative for 30-day projection
- `horizon60` : qualitative narrative for 60-day projection
- `horizon90` : qualitative narrative for 90-day projection
- `executiveWarning` : composite warning statement combining state, headcount, and market context

The engine uses ranges and language, not formulas. The exposure band is derived from the intersection of financial anchors and organisational state. When no financial data is provided but the state is DISORDERED, the band is "critical." When the state is MISALIGNED, the band is "high." When the state is DRIFTING, the band is "moderate." The system does not infer financial exposure from non-financial input.

The horizon narratives are state-specific templates that describe structural consequences in concrete terms. They do not predict specific financial losses. They describe what happens to decision authority, stakeholder alignment, and organisational capability when decisions remain unresolved.

The COI Engine serves a distinct function from the Cost of Delay Engine in `lib/diagnostics/cost-of-delay-engine.ts`. The Cost of Delay Engine operates at the diagnostic level with weighted scoring (urgency, ownership, clarity, accountability, state). The COI Engine operates at the governance level with state-based narrative generation. Both are server-side only. Neither exposes internal thresholds or multipliers.

Chapter 18: Execution Failure Layer

The Execution Failure Layer predicts why execution will stall. It operates on a state-based failure mode assessment model rather than a probabilistic prediction model. The system does not predict whether execution will fail. It identifies the specific structural conditions that make failure likely and prescribes the required correction.

Failure Mode Detection. The system detects failure modes from three inputs: clarity score, coherence score, and authority type (DIRECT, PROXY, or UNCLEAR). The detection is deterministic and implemented in

`lib/constitution/failure.ts` :

- Clarity below 45: "Weak problem articulation" — the decision is not stated clearly enough to execute.
- Coherence below 50: "Narrative fragmentation" — the story the organisation tells itself about the decision is internally inconsistent.
- Authority UNCLEAR: "Ambiguous decision rights" — no one has been clearly assigned the authority to act.
- Clarity below 55 AND coherence below 55: "Strategic drift" — both clarity and coherence are degraded simultaneously, indicating systemic rather than localised failure.
- Clarity below 40: "Foundational misalignment" — the organisation does not agree on the basic terms of the decision.

Required Correction. For each detected failure mode, the system prescribes a specific correction. These are not suggestions. They are structural requirements. If "Ambiguous decision rights" is detected, the required correction is: assign a named decision owner with explicit authority before proceeding. If "Narrative fragmentation" is detected, the required correction is: reconcile the competing internal narratives before attempting execution.

Integration with Intelligence Spine. Failure mode assessments feed into the spine through the deterministic output layer. When failure modes are detected, they inform the signal, the contradiction set, and the blocker class. They also influence the Decision Directive system, which may restrict or block escalation to higher-order surfaces when structural failure conditions are active.

The Execution Failure Layer is what separates this system from advisory models that assume execution will follow from good advice. This system assumes execution will fail unless the structural conditions for success are established first.

Chapter 19: Decision Memory v1

Decision Memory v1 is a persistence layer that stores public-safe decision summaries across sessions. It enables cross-session trend analysis without exposing internal scoring data.

Prisma Model: DecisionMemory

The DecisionMemory model stores:

- `userId` — optional, for authenticated users
- `organisationId` — optional, for enterprise contexts
- `sessionId` — optional, for anonymous session linking
- `source` — which product generated the memory (fast_diagnostic, constitutional, strategy_room, etc.)
- `state` — the organisational state at the time of the decision (ORDERED, DRIFTING, MISALIGNED, DISORDERED)
- `headline` — a one-line summary of the decision condition
- `summary` — a public-safe narrative summary
- `directive` — the recommended action
- `recommendations` — structured JSON array of specific recommendations
- `publicSignals` — optional JSON containing signals safe for client-facing output
- `escalationLabel` — optional label for escalation state
- `escalationLevel` — optional escalation level

Cross-Session Trend Analysis. When a user or organisation has multiple DecisionMemory records, the system can derive:

- **Dominant State** — what state appears most frequently across sessions. If MISALIGNED appears in 4 of 6 sessions, misalignment is the dominant condition, not an anomaly.
- **Repeated Conditions** — whether the same condition class (authority, definition, execution, instability) appears across multiple decisions. Repeated authority conditions suggest a structural governance problem, not a one-off decision failure.
- **Escalation Trend** — whether escalation levels are increasing, stable, or decreasing over time. Increasing escalation across sessions indicates the organisation's decision-making capacity is degrading.

- **Executive Summary** — a synthesised view across all stored memories, suitable for board-level reporting.

Decision Memory v1 is the foundation of the Enterprise layer's portfolio view. By persisting public-safe summaries, the system can show an organisation its decision-making pattern over time without exposing the proprietary scoring that produced each individual diagnosis. The records are indexed by `userId`, `organisationId`, `sessionId`, and `createdAt` for efficient retrieval and aggregation.

PART IV — SALES & POSITIONING

Chapter 20: Sales Doctrine

The sales doctrine of Abraham of London follows a single principle: authority before features.

The sales conversation does not begin with what the product does. It begins with a question the prospect cannot answer confidently. "How long has this decision been unresolved?" "Who owns it?" "What has it cost you so far?" If the prospect can answer these questions clearly and without hesitation, they do not need the system. If they cannot — and they almost never can — the system has already demonstrated its value before any feature has been described.

Pricing Structure:

- **Fast Diagnostic** — Free. Entry point. Under 2 minutes. Server-side scoring only.
- **Purpose Alignment Assessment** — Free. Personal diagnostic. Establishes mandate clarity.
- **Constitutional Assessment** — Included in engagement pathways. Structural diagnostic.
- **Executive Reporting** — **£295**. The first paid product. Delivers a formatted executive report with contradiction analysis, economic framing, and governance recommendations.
- **Strategy Room** — **£750-£1,250**. The enforcement product. 48-hour action-forcing session with locked decision, named owner, sequenced moves, and execution tracking. Pricing varies by complexity and organisational scale.
- **Retainers** — **£5,000-£50,000/month**. Ongoing decision authority infrastructure. Includes recurring diagnostics, multi-stakeholder deployment, longitudinal tracking, boardroom mode access, and portfolio-level decision intelligence. Retainer contracts are managed through the enterprise retainer system (`20260423_add_enterprise_retainer_contracts` migration).

Sales Process:

1. The prospect encounters the system through content (LinkedIn, canon, toolkits) or referral.
2. They run the Free Fast Diagnostic. The system produces an accurate reading.
3. Accuracy creates authority. The prospect recognises that the system has described their situation correctly.
4. The system offers the next surface (Executive Reporting or Strategy Room) with pricing anchored to the cost of inaction, not the cost of the product.
5. At no point does the sales process require the system to prove its methodology. The diagnostic output is the proof.

The Do-Not-Sell flag exists for prospects who should not proceed. If the system detects that the prospect lacks authority, has a trivial problem, or is seeking validation rather than resolution, the engagement is terminated. Revenue from wrong clients degrades the system's accuracy metrics and institutional credibility.

Chapter 21: Enterprise Case System

The Enterprise Case System produces evidence for institutional credibility without compromising client confidentiality. It operates on three principles:

Anonymised evidence only. No client is named. No organisation is identified. Cases are described in terms of condition class, contradiction pattern, intervention deployed, and verified outcome. The anonymisation is structural, not cosmetic — the case descriptions are written such that they could apply to any organisation facing that condition class.

No client naming. This is an absolute prohibition, not a preference. The system's credibility derives from accuracy, not from name-dropping. The statement "a FTSE 250 company in financial services" is never used because it narrows identification. The statement "an organisation experiencing authority-class contradiction with critical exposure" is used because it describes the condition without identifying the client.

Modelled outcomes. Case studies present modelled outcomes based on the system's diagnostic framework, not claimed client results. The system can state: "In cases where authority-class contradiction is resolved through Strategy Room intervention, the observed pattern is decision resolution within 14 days and structural change verification within 90 days." This is a statement about the system's observed patterns, not a promise about any specific client's results.

The Enterprise Case System feeds the Social Proof Engine (`lib/proof/social-proof-engine.ts`) which maintains observed outcome metrics and accuracy statistics. These are used in Boardroom Mode dossiers, Executive Reports, and sales materials. The evidence is drawn from the Evidence Ledger and verified outcomes, ensuring that every proof statement can be traced back to an actual verified impact record.

PART V — CONTENT & DISTRIBUTION

Chapter 22: Content Amplification

Content amplification operates on a 7-asset rotation model distributed primarily through LinkedIn, supported by the Abraham of London canon, glossary, toolkits, and frameworks.

The 7-Asset Rotation:

1. **Doctrine posts** — statements of principle from the founding thesis. These establish institutional authority. They do not sell. They define the category.
2. **Contradiction cases** — anonymised examples of specific contradiction patterns the system detects. These demonstrate diagnostic capability without revealing methodology.
3. **Economic framing pieces** — content that reframes the cost of inaction. These shift the prospect's mental accounting from "cost of intervention" to "cost of doing nothing."
4. **System capability posts** — descriptions of what the system does, focused on outputs not inputs. What the user sees, not how the system works.
5. **Authority challenge posts** — posts that challenge conventional decision-making assumptions. These repel wrong clients and attract right ones.
6. **Evidence posts** — anonymised outcome data from the system. Observed patterns, accuracy rates, resolution timelines.
7. **Engagement invitations** — direct calls to action. Run the diagnostic. See what it finds. Free. Under two minutes.

LinkedIn Campaign. The primary distribution channel is LinkedIn with a 21-post market-dominant entry campaign. This campaign was designed to establish category authority in the first wave. Each post follows a specific structural template: opening contradiction, institutional framing, system capability reference, and call to action. The posts are sequenced to build from doctrine (establishing the problem) through capability (establishing the system) to invitation (driving engagement). All 21 posts have been rewritten for market-dominant entry positioning.

Canon and Supporting Assets. The Abraham of London canon (`lib/canon.ts`) provides the foundational vocabulary. The glossary defines key terms. Toolkits provide frameworks that demonstrate the system's thinking without revealing its machinery. All supporting assets are public by design — they establish authority. The proprietary elements are the scoring, classification, and synthesis engines that sit behind the public face.

Chapter 23: Engagement Lane Architecture

The Engagement Lane Architecture organises the system's public surface into distinct pathways for different audience types. Each lane is a dedicated page with lane-specific content, messaging, and calls to action.

The 6 Lane Pages:

1. **Institutional** (/institutional) — for boards, governance systems, and institutional design work. Access tier: architect. This lane speaks to decision-makers responsible for organisational decision-making infrastructure. Messaging focuses on structural governance, board-level diagnostics, and multi-stakeholder deployment.
2. **Private Clients** (/private-clients) — for private engagement pathways, retainers, and advisory streams. Access tier: client. This lane speaks to individuals or small leadership teams facing specific high-stakes decisions. Messaging focuses on personal diagnostic depth, Strategy Room access, and ongoing retainer relationships.
3. **Education & Research** (/education-research) — for research notes, teaching editions, reading lists, and workshops. Access tier: public. This lane speaks to academics, students, and practitioners interested in decision authority as a field of study. Messaging focuses on intellectual contribution, research access, and workshop participation.
4. **Media** (/media) — for public commentary, interviews, press assets, and high-signal briefs. Access tier: public. This lane speaks to journalists, podcasters, and media professionals. Messaging focuses on expert commentary availability, press asset access, and briefing requests.
5. **Consulting** — for firms and practitioners who want to integrate decision authority methodology into their own practice. Messaging focuses on licensing, methodology access, and integration support.
6. **Contact** — the universal contact surface with lane-aware routing.

Lane-Aware Contact Form. The contact form accepts a `?type` query parameter that pre-selects the engagement lane. A visitor arriving from the Institutional page has `?type=institutional` pre-filled, routing their inquiry to the appropriate handling path. This ensures inquiries are categorised before they arrive, reducing response friction and enabling lane-specific follow-up sequences.

Homepage Strip. The homepage includes a navigation strip: "Not here to run a diagnostic?" This strip directs visitors who are not seeking a diagnostic to the appropriate engagement lane — media inquiries, research access, institutional engagement, or private consultation. The strip prevents the diagnostic from being the only entry point and ensures that non-diagnostic visitors find their lane without friction.

PART VI — SYSTEM ARCHITECTURE

Chapter 24: The 13-Layer Technical Architecture

The system is organised into 13 layers, each with a specific function. Layers are additive — higher layers depend on lower layers, but lower layers operate independently.

Layer 1: DETECTION

The detection layer identifies the decision condition from user input.

CaseObject. Created by the Fast Diagnostic. Contains the decision in the user's own words, claimed ownership, prior attempts, cost of delay, blockers, and forced action scenarios. Typed in `lib/decision/case-object.ts`. Classified into one of four condition classes: authority, definition, execution, instability.

C3 Fidelity Scorer. Scores input on three dimensions: Clarity (is the decision itself clear?), Context (is there enough surrounding information?), and Consequence (is the cost/stakes articulated?). Each dimension scores 0-1. The composite specificity score determines the processing tier. Implemented in `lib/decision/c3-fidelity-scorer.ts`.

Synthesis Engine with Server-Only Guard. The GovernedSynthesis output is LLM-generated analysis that is only produced when the C3 tier reaches FULL_SYNTHESIS (≥ 0.7). The synthesis engine is guarded by a server-only constraint — synthesis computation never occurs on the client. The `spine-guard.ts` enforces spine invariants.

Arbiter Tournament (5 rules). Every synthesis output must pass five mandatory validation rules before reaching the user. Implemented in `lib/decision/arbiter-tournament.ts`:

1. **Condition Integrity** — synthesis condition class must match deterministic class.
2. **Contradiction Alignment** — synthesis must reference the deterministic contradiction set.
3. **Move Validity** — recommended action must reference the user's specific blocker, not generic advice.
4. **Cost Consistency** — no invented costs without user-stated financial anchor.
5. **Avoidance Proof** — must reference the user's forcedAction or priorAttempt mismatch.

Hard violations reject the synthesis entirely. The user sees the mismatch. This is a market differentiator.

Layer 2: ENFORCEMENT

The enforcement layer ensures decisions are acted upon.

Intelligence Spine. The single persistent contract across all stages. Typed in `lib/decision/intelligence-spine.ts`. Never resets. Accumulates intelligence across stages.

Pressure Loop (48h/7d/14d). Time-based pressure escalation. The pressureIndex on the spine is monotonically increasing. Each missed checkpoint increases pressure. The pressure loop operates at three intervals: 48 hours (initial follow-up), 7 days (escalation), 14 days (final enforcement).

Breach Ladder. When committed actions are not taken, the breach is recorded on the spine (`execution.breach` , `execution.breachCount`). The Breach Ladder applies graduated consequences: first breach triggers re-engagement, second breach triggers escalation, repeated breaches trigger the Do-Not-Sell flag.

Pre-Commitment Gate. Captured during the Fast Diagnostic. The user explicitly states whether they are willing to act within 48 hours. This is not a question about interest. It is a commitment that the system tracks.

Layer 3: ANTI-MANIPULATION

The anti-manipulation layer protects the system's accuracy from gaming, flattery, and strategic misrepresentation.

Integrity Score. A 0-1 score on the spine (`integrityScore`). Below 0.5 indicates degraded integrity — the user's answers show patterns consistent with gaming or strategic misrepresentation. The score informs the system's confidence in its own output.

False Authority Detection. Flags discrepancies between claimed ownership and behavioural signals. When a user claims they own a decision but their answers reveal they cannot compel action, the `falseAuthority` flag is raised. Implemented through the CaseObject's comparison of `claimedOwner` against `actualOwnerCandidate` .

Economic Sanity. Detects implausible economic claims. If the stated cost of the decision is inconsistent with the described scope, the `economicSanitySuspicious` flag is raised. The system does not use implausible anchors for economic framing.

Phrasing Variation for Anti-Pattern-Detection. The system varies the phrasing of diagnostic questions across sessions to prevent users from learning the "right" answers. Consistent patterns across varied phrasings increase confidence. Inconsistent patterns across consistent phrasings raise integrity concerns.

Layer 4: SELF-GOVERNANCE

The self-governance layer ensures the system maintains its own standards.

North Star Metrics. The system tracks its own accuracy and impact metrics. These include: diagnostic accuracy (confirmed by user feedback), prediction accuracy (verified outcomes matching forecasts), resolution rate (proportion of diagnosed conditions that reach verified structural change), and authority metrics (proportion of engagements where the system's authority was confirmed by the client's actions).

Quality Gate (8-step validation). Every output passes through an 8-step validation process before reaching the user. The steps are documented in Appendix E. This gate is automated and enforced at the build level. No output bypasses the gate.

Layer 5: ECONOMIC PRESSURE

The economic pressure layer converts diagnostic findings into financial consequence language.

Cost Anchor. The user-provided financial anchor that grounds all economic projections. Without a cost anchor, the system produces qualitative exposure language only.

Cost of Inaction Engine. Server-side only (`lib/server/decision/cost-of-inaction.server.ts`). Produces exposure bands and horizon narratives from state and financial input. See Chapter 17 for full documentation.

Do-Not-Sell. A flag that terminates engagement when the system determines that selling to this client would degrade institutional credibility. Triggered by: trivial problems, absent authority, validation-seeking behaviour, or economic unsuitability.

Time Decay. The mathematical model for option compression over time. Implemented through the `DefaultPathForecast`'s `optionDecayRate` (options lost per month, 0-1) and `controlShiftProbability` (probability that decision control transfers away from stated owner, 0-1). Base decay rates vary by condition class: authority 0.3, execution 0.4, definition 0.2, instability 0.15.

Layer 6: EVIDENCE PERMANENCE

The evidence permanence layer creates an immutable record of decisions, actions, and outcomes.

Evidence Ledger. An append-only record implemented in `lib/constitution/evidence-ledger.ts` . Every stage appends. Nothing is overwritten. Each entry records: spine ID, stage, contradiction, action required, action taken, verified impact, pressure index, integrity score, condition class, and cost estimate. The ledger is built from the spine history and can only be appended to, never modified.

Decision Credit Score. A derived score based on the user's historical pattern of making commitments and following through. Users who consistently commit and execute build decision credit. Users who consistently commit and breach degrade their score. This score influences future system recommendations.

Contradiction Convergence. Cross-session analysis of contradiction patterns. When the same contradiction appears across multiple sessions or multiple stakeholders, it is flagged as convergent — indicating a structural condition rather than a situational one.

Layer 7: SOCIAL PROOF

The social proof layer provides evidence of the system's effectiveness.

Observed Outcomes. Aggregated, anonymised data from the Outcome Verification loop. The system can report: "In X% of cases where condition Y was detected and intervention Z was applied, verified structural change occurred within N days."

Accuracy Metrics. Aggregated accuracy feedback from users. The system tracks what proportion of users confirmed the diagnosis as accurate (yes), partially accurate (partial), or inaccurate (no).

Anonymised Case Studies. See Chapter 21. Evidence drawn from the Evidence Ledger and verified outcomes. No client naming. Modelled outcomes only.

Layer 8: BOARDROOM MODE

The boardroom layer produces governance-grade decision documents.

Board Dossier. A multi-section document generated from the Intelligence Spine. Sections use factual, confrontational, and quantified tones. Includes objection handling and decision path analysis. See Chapter 14.

Multi-Stakeholder Collision Detection. Identifies where multiple stakeholders' inputs diverge on the same decision. Surfaces authority gaps, cost divergence, and narrative inconsistencies across respondents.

Layer 9: MULTI-ACTOR TRUTH

The multi-actor truth layer extends collision detection to organisational-level analysis.

Collision Detection. Implemented in `lib/constitution/multi-user-collision.ts`. When multiple users from the same organisation run diagnostics on related decisions, the system detects where their answers collide — where they disagree about ownership, urgency, cost, or blockers.

Authority Gaps. Situations where no stakeholder claims ownership, or where multiple stakeholders claim ownership of the same decision. Both conditions are structural failures that the system surfaces.

Cost Divergence. Situations where different stakeholders assign radically different economic values to the same decision. This divergence is itself a diagnostic signal — it indicates that the organisation has not agreed on what the decision is worth, which means they have not agreed on what it is.

Layer 10: INBOUND FILTERING

The inbound filtering layer routes incoming engagement to the appropriate surface.

Tier Routing. The access tier system (public, member, inner_circle, restricted, client, legacy, architect, owner, top_secret) determines which surfaces a user can access. Implemented through `lib/access/tier-policy.ts` and `lib/access/tier-map.ts`.

Strategy Room Gate. The Decision Directive system evaluates whether a user's diagnostic state qualifies them for Strategy Room access. If structural failure is detected, access is blocked. If a mandate vacuum exists, access is restricted. See Chapter 11.

Inbound Filter. The contact form and engagement pathways use lane-aware routing (`?type` parameter) to direct inquiries to the appropriate handling path. See Chapter 23.

Layer 11: ESCALATION ENGINE

The escalation engine applies time-based pressure to unresolved decisions.

Escalation Timeline: 24h, 48h, 72h, 5d, 7d. The system applies escalating pressure at five intervals:

- **24h** — initial follow-up. Confirms the decision is still active. Restates the diagnosed contradiction.
- **48h** — commitment checkpoint. Evaluates whether the Pre-Commitment Gate promise has been honoured. If not, pressure increases.
- **72h** — escalation warning. The system warns that the cost of inaction is now higher than it was at diagnosis.

- **5d** — breach detection. If no action has been taken, the system records a breach and applies the Breach Ladder.
- **7d** — final enforcement. The system delivers the 7-day consequence projection and either escalates to a higher surface or triggers the Do-Not-Sell flag.

The auto-escalation sweep (`lib/constitution/escalation-engine.ts`) runs periodically, checking open interventions. Cases with seriousness above 75 and confidence above 0.6 are automatically escalated to the Strategy Room.

Layer 12: RETARGETING

The retargeting layer re-engages users who have received a diagnosis but have not acted.

Pressure Loop. A time-based re-engagement sequence that references the user's specific diagnostic findings. The loop does not send generic follow-ups. It restates the contradiction, updates the cost of inaction projection, and asks whether the condition has changed.

Follow-Up System. Managed through `lib/follow-up/retargeting-segments.ts` . Segments users by diagnostic severity, engagement level, and response history. Follow-up content is personalised to the user's condition class and exposure band.

Layer 13: ENTERPRISE

The enterprise layer provides portfolio-level decision intelligence.

Decision Memory. See Chapter 19. Cross-session persistence of public-safe decision summaries.

Cross-Session Signal Aggregation. The system aggregates signals across multiple sessions for the same user or organisation. This enables trend detection: is the organisation's decision-making capacity improving, stable, or degrading?

Portfolio View. For enterprise retainer clients, the system provides a portfolio-level view of all active decisions, their condition classes, exposure bands, escalation states, and resolution status. This is the executive dashboard for decision authority infrastructure. Implemented through `lib/enterprise/enterprise-portfolio.ts` .

Chapter 25: Delivery Model & Operating Cadence

Technology Stack. The system is built on Next.js with TypeScript. The database is PostgreSQL accessed through Prisma ORM. PDF generation is handled through a dedicated rendering pipeline. The system deploys on Netlify with serverless functions for API routes.

Deployment Cadence. Changes are deployed through a CI/CD pipeline. Every deployment passes through the Quality Gate (Appendix E). No deployment bypasses the gate. The deployment model is continuous — there are no scheduled release windows.

Data Architecture. The database schema is managed through Prisma migrations. As of April 2026, the migration history includes: access tier expansion, proof evidence, strategy room tables, team assessment campaigns, diagnostic evidence graph, purpose alignment canonical results, longitudinal/multi/outcome infrastructure, strategy room

enforcement persistence, enterprise retainer contracts, AI adaptive decision authority, enterprise-capable foundation, PAL advantage fields, verification token, and calibration loop.

API Structure. The API is organised by domain:

- `/api/diagnostics/` — diagnostic submission, re-entry, outcome verification, longitudinal, multi-stakeholder
- `/api/constitutional/` — constitutional assessment, audit, export, appeal
- `/api/strategy-room/` — session management, execution tracking, conversion
- `/api/executive-reporting/` — entitlements, PDF export, boardroom PDF, intervention export
- `/api/purpose-alignment/` — assessment, report, reminders
- `/api/alignment/enterprise/` — enterprise campaigns, organisations, assessments
- `/api/admin/` — decision intelligence, commercial, campaign management, governance
- `/api/sovereign/` — authenticated session management

Operating Cadence. The system operates on two cadences:

1. **Diagnostic cadence** — real-time. Diagnostics are processed immediately upon submission. Scoring is server-side. Results are returned synchronously.
 2. **Enforcement cadence** — time-based. Escalation, breach detection, and follow-up operate on the 24h/48h/72h/5d/7d timeline. The auto-escalation sweep runs periodically to detect cases requiring automatic Strategy Room escalation.
-

PART VII — CAPITAL & INVESTOR RELATIONS

Chapter 26: Capital Narrative

Abraham of London creates a new category: **Decision Authority Infrastructure (DAI)**.

The capital narrative begins with a market observation: every organisation spends heavily on intelligence (consultants, analytics, data platforms) and on execution (project management, process automation, operational tools). But no one has built infrastructure for the gap between intelligence and execution — the decision itself.

Decisions are the bottleneck. Not information. Not execution capability. The decision. And decisions fail for structural reasons: unclear authority, undefined scope, unpriced consequences, and absent enforcement. These are infrastructure problems, not information problems.

DAI addresses this gap by providing the structural layer that forces decisions to resolution. It is not consulting (which depends on client willingness), not analytics (which provides information without enforcement), and not project management (which assumes the decision has already been made). It is the infrastructure that ensures decisions get made, tracked, and verified.

The market opportunity is defined by the size of the consulting and advisory market (\$300B+ globally) and the observation that a significant proportion of advisory engagements fail to produce action — not because the advice was wrong, but because the decision architecture was inadequate. DAI captures value from this failure mode.

The competitive advantage is structural: the system's intelligence accumulates across sessions (the Intelligence Spine never resets), its enforcement mechanisms create accountability that advisory cannot, and its verification loop measures actual impact rather than claiming it. The longer a client uses the system, the more accurate it becomes. The longer a competitor operates without these mechanisms, the further behind they fall.

Chapter 27: Investor Memo

The Problem. Organisations spend billions on consultants and analytics. Decisions still stall. The problem is not information — it is decision authority. No infrastructure exists to identify unresolved decisions, price their consequences, force resolution, and verify outcomes.

The Solution. Abraham of London is Decision Authority Infrastructure. The system identifies contradictions in decision-making through structured diagnostics, prices the cost of inaction through server-side economic engines, forces decisions through enforcement mechanisms (Pre-Commitment Gate, Strategy Room, Escalation Engine), and verifies outcomes through the Outcome Verification loop.

The Product. A tiered product system from free diagnostics (Fast Diagnostic, Purpose Alignment) through paid enforcement (Executive Reporting at £295, Strategy Room at £750-£1,250) to enterprise retainers (£5,000-£50,000/month). Each tier increases diagnostic depth, enforcement pressure, and accountability.

The Technology. A 13-layer architecture with server-side scoring, deterministic classification, LLM synthesis with hard validation (5-rule Arbiter Tournament), persistent Intelligence Spine, cross-session Decision Memory, and append-only Evidence Ledger. All proprietary logic is server-only. The public surface demonstrates capability without revealing machinery.

The Traction. The system is live with functional diagnostics, strategy room, executive reporting, enterprise campaigns, and boardroom mode. The 21-post LinkedIn campaign is deployed for market-dominant entry. The engagement lane architecture provides multi-audience reach.

The Ask. Seed capital to fund: first ten enterprise accounts, content amplification at scale, engineering depth for the 13-archetype taxonomy and NEGLIGIBLE/SYSTEMIC severity tiers, and sales infrastructure.

The Return. Revenue model is tiered: free-to-paid conversion from diagnostics to Executive Reporting (£295), upsell to Strategy Room (£750-£1,250), land-and-expand to retainers (£5K-£50K/month). Retainer revenue is recurring and deepens with usage. The Intelligence Spine creates structural lock-in: switching costs increase as decision history accumulates.

Chapter 28: Investor Call Script

Opening (2 minutes). "Every organisation you've invested in has had a decision that should have taken two weeks and took six months. Not because they lacked information. Because no one had the authority, the pressure, or the accountability to force it to conclusion. We've built infrastructure for that."

Problem Statement (3 minutes). Describe the decision authority gap. Use anonymised examples from the system's case library. Quantify the cost of indecision using the economic framing model. Establish that this is a structural problem, not an information problem.

Solution Overview (5 minutes). Walk through the four operations: identify contradiction, price consequence, force decision, verify outcome. Demonstrate the Fast Diagnostic live — share screen, run a diagnostic, show the output. Let the system speak for itself. Do not explain the methodology. Show the result.

Product and Revenue (5 minutes). Present the tiered model: free diagnostics as entry, Executive Reporting as first conversion, Strategy Room as enforcement product, retainers as recurring revenue. Show the pricing logic: price is always less than the cost of inaction the system has identified. The sale is not "buy this product." The sale is "you cannot afford not to resolve this decision."

Technology Moat (3 minutes). Explain the Intelligence Spine — the fact that it never resets, that every session adds intelligence, that the system becomes more accurate with use. Explain the Arbiter Tournament — the fact that the system rejects its own output when it fails validation. These are not features. They are structural advantages that cannot be replicated by adding a feature to an existing product.

Ask and Terms (2 minutes). State the ask. Present terms. Close with: "We are not asking you to believe in our methodology. We are asking you to verify it. Run the diagnostic. If the findings are accurate, the investment thesis follows."

Chapter 29: Objection Handling — Investor Warbook

The Investor Warbook contains 17 investor objections with structured responses.

1. **"This is just consulting with software."** No. Consulting depends on the consultant's continued involvement. This is infrastructure that operates independently. Consulting gives advice. This system enforces decisions.
2. **"The market is too small."** The global management consulting market is \$300B+. We are not competing with it. We are capturing the value it fails to deliver — the decisions that were correctly diagnosed but never resolved.
3. **"Why would anyone pay for this when they can hire McKinsey?"** McKinsey tells you what to do. We ensure it gets done. They are complementary, not competitive. McKinsey clients are our best prospects — they already have the diagnosis but lack the enforcement.
4. **"The technology isn't defensible."** The Intelligence Spine's cross-session accumulation creates data moats. The Arbiter Tournament's validation logic is proprietary. The Evidence Ledger's longitudinal records cannot be replicated without time. Defensibility increases with usage.
5. **"How do you acquire customers?"** Free diagnostic as entry point. Content amplification through LinkedIn. Lane-specific engagement architecture. The diagnostic is the sales tool — if it is accurate, the system sells itself.
6. **"What's the churn risk?"** The Intelligence Spine creates switching costs. Decision history does not transfer. The longer a client uses the system, the more valuable the accumulated intelligence becomes.
7. **"Can this scale without human intervention?"** The diagnostic layer is fully automated. The enforcement layer is time-triggered. The Strategy Room requires human facilitation but the surrounding infrastructure (escalation, tracking, breach detection) is automated.
8. **"Who is the competition?"** No one builds decision authority infrastructure. Consulting firms sell advice. Analytics platforms sell data. Project management tools track execution. No one tracks the decision itself. We are creating the category.
9. **"What if the diagnosis is wrong?"** The system measures its own accuracy through the Outcome Verification loop. Wrong diagnoses are detected, recorded, and fed back into the system. Self-correction is built in, not bolted on.
10. **"This sounds too aggressive for enterprise."** Enterprise does not need gentle. Enterprise needs accurate. The system's directness is its value proposition. Organisations do not change because someone asked nicely.
11. **"What's the LTV:CAC?"** Free diagnostic means near-zero CAC for entry. Executive Reporting at £295 covers acquisition cost. Strategy Room at £750-£1,250 is profit. Retainers at £5K-£50K/month are high-LTV recurring revenue.
12. **"How do you handle data privacy?"** All diagnostic data is processed server-side. Decision Memory stores public-safe summaries only. No internal scoring data is exposed. Client data is isolated by organisation.
13. **"What's the team?"** The system was designed by a single institutional architect with full doctrinal control. This is an advantage at this stage — no design-by-committee, no diluted vision.
14. **"Is the AI defensible?"** The AI (LLM synthesis) is the least important part of the system. The deterministic layer, the C3 scorer, the Arbiter Tournament, the Evidence Ledger — these are the defensible components. The AI is a synthesis layer that is validated, not trusted.
15. **"What if someone copies the methodology?"** The methodology is public. The machinery is proprietary. Copying the framework is easy. Replicating the Intelligence Spine, the Evidence Ledger, the cross-session memory, and the Arbiter Tournament requires building the entire system. See Chapter 35: Disclosure Principle.

16. **"Can you show me case studies?"** Anonymised evidence only. We do not name clients. We describe conditions, interventions, and verified outcomes. See Chapter 21.
17. **"What's the exit path?"** Enterprise SaaS acquisition by governance, risk, or consulting platforms seeking to add decision enforcement to their portfolio. Alternatively, category leadership in Decision Authority Infrastructure as a standalone public company.

Chapter 30: Fundraising Strategy

Stage: Pre-Seed / Seed. The system is functional. The products are live. The technology is built. The fundraise is for market entry and first accounts, not for product development.

Target Raise. Sufficient to fund 18 months of operations: first ten enterprise accounts, one sales hire, content amplification, and engineering depth for target features (13-archetype taxonomy, NEGLIGIBLE/SYSTEMIC severity tiers).

Investor Profile. Seeking investors who understand:

- Category creation (not competing in an existing market)
- Infrastructure businesses (not feature-based SaaS)
- Founder-led vision (not consensus-driven product development)
- Long-cycle enterprise sales (not immediate viral growth)

Fundraising Process:

1. **Warm introduction.** Every investor interaction begins with a warm introduction from someone who has experienced the system. Cold outreach is not used. The system's own doctrine applies: authority before features.
2. **Diagnostic demonstration.** Every investor call includes a live diagnostic. The investor experiences the system's accuracy firsthand. This is not a demo. It is a proof.
3. **Memo delivery.** The Investor Memo (Chapter 27) is delivered after the diagnostic, not before. The diagnostic creates the context for the memo. Without the diagnostic experience, the memo is just another pitch deck.
4. **Term negotiation.** Terms are standard for the stage. The non-negotiable term is doctrinal control — the founding vision is not subject to board override during the category creation phase.
5. **Close.** The close mirrors the system's own methodology: name the decision, name the owner, set the deadline, track the execution. The investor either acts or does not. The system does not chase.

PART VIII — OPERATIONAL DISCIPLINE & GOVERNANCE

Chapter 31: The First Ten Accounts

The first ten accounts define the institution. They are not just customers. They are evidence. Every subsequent sale, every case study, every investor conversation, and every content piece will reference the patterns established by these ten accounts.

Selection Criteria. The first ten accounts must satisfy all of the following:

1. **Real decision at stake.** The organisation must have a genuine unresolved decision with material consequence. Test cases, pilots, and "let's see what it does" engagements are rejected.
2. **Authority to act.** The primary contact must have the authority to force the decision to conclusion, or must have direct access to someone who does. Engagements with middle managers who need approval chains are deferred.
3. **Willingness to verify.** The organisation must agree to participate in the Outcome Verification loop. This is non-negotiable. Without verification, the engagement produces no evidence.
4. **Appropriate scale.** The decision must be significant enough to justify the system's intervention but not so complex that resolution within the system's timeframe is impossible.
5. **Condition class diversity.** The ten accounts should span all four condition classes (authority, definition, execution, instability) to build evidence breadth.

Engagement Protocol for First Ten:

- Run the full diagnostic ladder: Fast Diagnostic, Constitutional Assessment, and Strategy Room.
- Apply the enforcement model at full intensity. Do not soften for early clients.
- Capture Outcome Verification at 30, 60, and 90 days.
- Record every interaction in the Evidence Ledger.
- Produce anonymised case documentation for each account.

Success Metrics for First Ten:

- 80%+ diagnostic accuracy (confirmed by client feedback)
 - 70%+ action taken within committed timeframe
 - 50%+ verified structural change at 90 days
 - 100% Outcome Verification participation
 - Zero client naming in any external material
-

Chapter 32: What Not To Do

The following 24 prohibitions govern all institutional behaviour. Violations are not corrected. They are treated as breaches of institutional discipline.

1. Do not sell to organisations that lack decision authority.
 2. Do not soften diagnostic findings to maintain client comfort.
 3. Do not fabricate financial projections without user-provided anchors.
 4. Do not name clients in any external material, ever.
 5. Do not bypass the Arbiter Tournament for any reason.
 6. Do not deliver synthesis output that has failed hard validation.
 7. Do not reset the Intelligence Spine between stages.
 8. Do not decrease the pressureIndex under any circumstances.
 9. Do not offer discounts that undermine the pricing doctrine.
 10. Do not compete on warmth, accessibility, or ease of use.
 11. Do not describe the system as consulting, coaching, or advisory.
 12. Do not expose internal scoring thresholds or classification logic to clients.
 13. Do not deploy Boardroom Mode for non-board-level issues.
 14. Do not proceed with Strategy Room when the Decision Directive blocks it.
 15. Do not suppress the Do-Not-Sell flag for revenue purposes.
 16. Do not promise outcomes the system has not verified.
 17. Do not use client data for marketing without explicit authorisation.
 18. Do not skip the Quality Gate for any deployment.
 19. Do not add features that trade accuracy for user comfort.
 20. Do not accept investment that requires diluting doctrinal control.
 21. Do not publish content that reveals proprietary scoring or synthesis logic.
 22. Do not engage with media that requires simplifying the system to fit a headline.
 23. Do not hire salespeople who cannot explain the founding problem.
 24. Do not treat the first ten accounts as anything less than institutional evidence.
-

Chapter 33: Quality Assurance & Compliance

Quality Gate. Every release passes through the 8-step Quality Gate documented in Appendix E. No exceptions. No manual overrides. The gate is automated and enforced at the build level.

Diagnostic Quality. Every diagnostic output is validated against the following standards:

- CaseObject fields must be populated from actual user input, not defaults or inferences.
- C3 scores must be computed server-side using the documented scoring function.
- Deterministic output must include a valid condition class, signal, and contradiction set.

- Synthesis (when produced) must pass the 5-rule Arbiter Tournament.
- Economic projections must reference user-provided anchors or default to qualitative language.

Data Integrity. The Evidence Ledger is append-only. Historical entries are never modified. The spine's pressureIndex is monotonically increasing. The integrityScore is derived from behavioural signals, not manual assignment.

Compliance. Client data is processed server-side and stored in PostgreSQL. Decision Memory stores public-safe summaries only — no internal scoring data is persisted in client-facing records. The system does not store raw diagnostic answers after scoring is complete. Access tiers control visibility at the application level.

Audit Trail. The audit system (`lib/audit/audit-logger.server.ts`) records significant system actions including diagnostic submissions, report generation, strategy room sessions, and escalation events. Audit logs are append-only and timestamped.

Monitoring. The system includes observability infrastructure (`lib/constitution/observability-dashboard.ts` , `observability-store.ts`) for tracking system health, performance metrics, and operational anomalies.

Chapter 34: Engagement Termination

Engagements are terminated under three conditions:

Condition 1: Do-Not-Sell Triggered. When the system detects that engagement would be unproductive — because the client lacks authority, the problem is trivial, or the client is seeking validation rather than resolution — the `doNotSellTriggered` flag is raised on the spine. Engagement is terminated. The client is informed that the system has determined this is not a productive engagement and no further services will be offered.

Condition 2: Repeated Breach Without Action. When the client repeatedly commits to action and repeatedly fails to execute, and the system has exhausted its escalation ladder, the engagement is terminated. The system cannot force action. It can only create the conditions under which action becomes the path of least resistance. When those conditions fail, the engagement has reached its structural limit.

Condition 3: Integrity Degradation. When the integrityScore drops below 0.3, indicating persistent gaming, strategic misrepresentation, or deliberate manipulation of the diagnostic process, the engagement is terminated. The system's authority depends on accurate input. When input integrity is compromised, diagnostic accuracy is compromised, and the system's core value proposition is void.

Termination Protocol:

1. The system generates a final diagnostic summary documenting the condition at termination.
2. The summary is appended to the Evidence Ledger with a termination flag.
3. The client receives a termination notice that includes the reason category but not the internal scoring that triggered it.
4. All client data is retained in the system for institutional learning but is flagged as terminated.
5. Re-engagement is possible after a minimum 90-day cooling period, subject to a fresh diagnostic that demonstrates changed conditions.

Termination is not punitive. It is structural. The system recognises the limits of its own intervention model and withdraws rather than degrading its output quality to maintain revenue.

PART IX — IP PROTECTION & COMPETITIVE DEFENSE

Chapter 35: Disclosure Principle

The disclosure principle governs what the system reveals publicly and what remains proprietary:

"The user sees what it does. The competitor cannot infer how."

This principle operates at every layer of the system. The user receives a diagnostic result that includes a condition class, contradiction set, economic framing, and concrete move recommendation. This output is visible, understandable, and valuable. But the output does not reveal how the condition class was derived, how the contradiction was detected, how the economic framing was computed, or how the recommended move was validated.

The disclosure principle is not security through obscurity. The system's foundational concepts — condition classes, contradiction detection, economic consequence, enforcement — are public. They are published in the canon, discussed in content, and described in the engagement lane pages. This public disclosure serves authority: it demonstrates that the institution has a coherent intellectual framework.

What remains proprietary is the machinery:

- The C3 Fidelity Scorer's weighting and threshold logic
- The Arbiter Tournament's five validation rules and their implementation
- The scoring functions in `lib/diagnostics/scoring.ts`
- The cost computation logic in `lib/server/decision/cost-of-inaction.server.ts`
- The failure mode detection thresholds in `lib/constitution/failure.ts`
- The classification boundaries in `lib/constitution/classification.ts`
- The synthesis governance and rejection logic
- The integrity score derivation
- The pressure index computation

A competitor reading this manual will understand what the system does and why. They will not understand how to replicate it. That asymmetry is by design.

Chapter 36: Public vs Proprietary Layers

The system is divided into two layers with a clear boundary between them.

Public Layer — Authority Assets:

The following are public by design. They establish the institution's authority and market position:

- **Canon** (`lib/canon.ts`): The foundational vocabulary and conceptual framework. Published openly. Anyone can read and reference the canon. This builds category awareness.
- **Glossary**: Definitions of key terms (condition class, contradiction, decision authority, enforcement, etc.). Published on the website and in content.
- **Foundations**: The founding problem, the institutional thesis, the four operations. Published in the LinkedIn campaign, the engagement lanes, and this manual's doctrinal chapters.
- **Toolkits and Frameworks**: Structured thinking tools that demonstrate the system's approach. These help potential clients understand the framework without accessing the machinery.

The public layer is generous. It gives away concepts freely. This is not leakage. It is strategy. When the system's concepts become the industry vocabulary — when people say "decision authority gap" and "condition class" — the institution that defined those terms owns the category.

Proprietary Layer — Server-Only Machinery:

The following are server-only and never exposed to the client or the public:

- **Scoring logic**: All scoring functions, weighting coefficients, and threshold boundaries. The C3 Fidelity Scorer, the severity scoring function, the readiness classifier, and the posture classifier.
- **Classification logic**: The exact boundaries between condition classes, severity levels, and score bands. The system reports the classification. It does not reveal where the boundary lies.
- **Synthesis governance**: The Arbiter Tournament rules, the synthesis engine, and the rejection logic. The system reports that synthesis was accepted or rejected. It does not reveal why.
- **Economic computation**: The Cost of Inaction Engine, the Cost of Delay Engine, and the Default Path Forecast derivation logic. The system reports exposure bands and horizon narratives. It does not reveal the computation.
- **Integrity and pressure mechanics**: The integrityScore derivation, the pressureIndex computation, and the breach detection logic. The system applies these. It does not explain them.

Controlled Ambiguity. All public documentation maintains controlled ambiguity about the proprietary layer. The system describes what it does ("identifies contradictions," "prices consequences," "forces decisions") without describing how. This is not deception. It is discipline. The user does not need to know how the engine works. They need to know that it works. The proof is in the diagnostic output's accuracy, not in the code's documentation.

The boundary between public and proprietary is enforced technically through the `server-only` import guard on sensitive modules and architecturally through the separation of client-facing API responses (which contain classifications and narratives) from internal scoring objects (which contain raw scores, weights, and thresholds).

APPENDICES

Appendix A: Pricing Matrix

Product	Price	Type	Access
Fast Diagnostic	Free	Entry-level diagnostic	Public
Purpose Alignment Assessment	Free	Personal diagnostic	Public
Constitutional Assessment	Included in engagement	Structural diagnostic	Member+
Executive Reporting	£295	Formatted executive report	Client
Strategy Room	£750 - £1,250	Enforcement session (48h)	Client
Boardroom Mode	Included in Enterprise tier	Board dossier generation	Architect
Enterprise Retainer (Tier 1)	£5,000/month	Recurring diagnostics + tracking	Client
Enterprise Retainer (Tier 2)	£15,000/month	Multi-stakeholder + portfolio	Architect
Enterprise Retainer (Tier 3)	£50,000/month	Full institutional deployment	Owner

All pricing is in GBP. Enterprise retainer pricing varies by organisational scale, number of active decisions tracked, and depth of multi-stakeholder deployment. The pricing principle is constant: the price is always less than the cost of inaction the system has identified.

Appendix B: Severity Taxonomy

The system currently implements a 4-tier severity taxonomy:

Severity	Score Range	Verdict
LOW	80-100%	Stable with minor correction points
MODERATE	60-79%	Recoverable, but drift is present
HIGH	40-59%	Material weakness detected
CRITICAL	0-39%	Immediate correction required

The DiagnosticSeverity type in the codebase defines four values: "low", "moderate", "high", "critical".

TARGET — not yet implemented: The target taxonomy adds two additional tiers:

Severity	Description
NEGLIGIBLE	No material decision risk detected. Monitoring only.
SYSTEMIC	Condition affects multiple interconnected decisions. Organisational-level intervention required.

The NEGLIGIBLE tier would sit above LOW, capturing cases where no intervention is warranted. The SYSTEMIC tier would sit above CRITICAL, capturing cases where the failure is not in a single decision but in the organisation's decision-making system itself. Both tiers require additional detection logic and scoring thresholds that have not yet been implemented.

Appendix C: Condition Class Taxonomy

The system implements 4 condition classes, defined as the `ConditionClass` type in `lib/decision/case-object.ts` :

Condition Class	Definition	Base Decay Rate
Authority	The decision has no legitimate owner, or the stated owner lacks actual authority to compel action.	0.30
Definition	The decision itself is not clearly defined. Stakeholders disagree on what the question is.	0.20
Execution	The decision has been made but is not being executed. Structural blockers prevent action.	0.40
Instability	The conditions around the decision are shifting faster than the organisation can adapt.	0.15

Base decay rates represent options lost per month (0-1 scale) and are defined in the DefaultPathForecast derivation logic. Actual decay rates are adjusted upward based on user-provided specificity in cost-of-delay input and the presence of forced action vs. blocker discrepancies.

The condition class is the primary classification axis for every case. It determines the deterministic output, the default path forecast, the economic framing tone, and the recommended intervention pathway.

Appendix D: 13 Contradiction Archetype Registry

STATUS: TARGET — not yet implemented. The current system detects contradictions through rule-based pattern matching with 4 detection rules. The target is a taxonomy of 13 named archetypes distributed across the 4 condition classes.

Authority Class (4 archetypes)

1. **The Phantom Owner** — someone is named as decision owner but has no mechanism to compel action. The title exists. The authority does not.
2. **The Consensus Trap** — the organisation requires consensus before deciding, ensuring no individual can be held accountable for the decision.
3. **The Proxy Chain** — authority is delegated through so many layers that no one knows who actually decides.
4. **The Veto Ghost** — someone who is not the named owner can block the decision without being visible in the authority structure.

Definition Class (3 archetypes)

5. **The Moving Target** — the decision has been redefined so many times that its original urgency has been diluted.
6. **The Scope Creep Mask** — what began as a specific decision has expanded to encompass everything, ensuring it encompasses nothing actionable.
7. **The False Binary** — the decision is framed as two options when the real issue is that neither option addresses the actual problem.

Execution Class (3 archetypes)

8. **The Pilot Purgatory** — the decision was "made" but implementation is perpetually in pilot, proof-of-concept, or phase one.
9. **The Resource Mirage** — resources are nominally committed but practically unavailable, ensuring the decision cannot be executed.
10. **The Partial Execute** — enough action is taken to create the appearance of progress while the core issue remains unresolved.

Instability Class (3 archetypes)

11. **The Stable Decision in Unstable Ground** — the organisation has made a clear decision but the conditions it depends on are shifting, making the decision progressively less relevant.
12. **The Adaptive Freeze** — the environment is changing but the organisation has frozen its decision-making, hoping stability will return.
13. **The Reactive Loop** — the organisation is making rapid decisions in response to each change, but the decisions contradict each other because there is no stable reference point.

Each archetype, when implemented, will have: a detection signature (patterns in CaseObject fields that indicate its presence), a scoring weight (how much it contributes to severity), and a prescribed intervention (the specific system surface best suited to address it).

Appendix E: Quality Gate Checklist

The 8-step Quality Gate is enforced on every deployment. No step is optional. No manual override exists.

Step	Gate	Description	Fail Action
1	Prisma	Prisma schema validates. Migrations are clean. No drift between schema and database.	Block deployment.
2	TypeScript	Full TypeScript compilation with strict mode. Zero type errors.	Block deployment.
3	PDF	PDF generation pipeline produces valid output for all registered templates. No rendering errors.	Block deployment.
4	MDX Integrity	All MDX content files parse without errors. Frontmatter validates against schema.	Block deployment.
5	MDX Gate	MDX content gate verifies that no content references undefined components, broken links, or missing assets.	Block deployment.
6	Unit Tests	All unit tests pass. Test coverage thresholds are met. No skipped tests without documented justification.	Block deployment.
7	Build	Full production build completes without warnings or errors. Bundle size is within acceptable limits.	Block deployment.
8	Integration	End-to-end validation of critical paths: diagnostic submission, scoring, result rendering, PDF export, and API response contracts.	Block deployment.

Every step produces a pass/fail result. Any fail blocks the deployment. The gate is not a suggestion. It is a structural enforcement mechanism that prevents degraded output from reaching production. The gate protects the institution's most important asset: accuracy.

End of Manual.

Abraham of London — Decision Authority Infrastructure Version 2.0 — April 2026 Classification: Internal — Institutional